

**PA02: Mandelbrot Computations**  
**Complex, Independent Computations Sequentially**  
**vs. Static and Dynamic Parallel Modalities**

**Terence Henriod**  
**Dr. Fred Harris**  
**CS651: Parallel Computing**  
**Monday March 10, 2014**

## Introduction

Some tasks require many complex computations. Often such computations are independent of one another, so performing the computations one at a time, one after another, is less than ideal. If the computations can be performed alongside one another, i.e. in parallel, we can perform the computations in amounts of time that are more reasonable to humans. It should also be noted that different parallel computation modalities exist, two of which are explored here: static and dynamic job allocation.

## Theory

### *Mandelbrot Set*

In this report, we use the example of the Mandelbrot set. The Mandelbrot set is a set of points that comprise a boundary that creates a fractal image. This image is seen in pop-math, and requires (relatively) extensive computation to produce each point (or pixel) in the set (image). In short, each Mandelbrot point/pixel color/shade is found by computing the following formula iteratively, using the result in the previous iteration to prime the next:

$$z_{k+1} = z_k^2 + c \quad [1]$$

where  $z_k$  is the  $k$ th iteration of the formula and  $c$  is the complex value representing the location of the pixel in the complex plane. The computation stops after either a predetermined number of iterations or when it can be seen that the formula will diverge, that is, if the formula will produce a sequence of zeroes or progress to  $\pm\infty$ . The computation and theory of the Mandelbrot set is not the focus of this report, so it will not be discussed further.

### *Speedup*

When computing things in parallel, it is advantageous to know how much faster the parallel algorithm/program runs compared to a sequential one. This is one metric used to evaluate the quality of a parallel algorithm, the *speedup factor*. The speedup factor represents how many times faster the parallel algorithm using  $p$  processors is compared to the sequential one. The speedup factor,  $S(p)$ , is computed as follows:

$$S(p) = \frac{t_s}{t_p} = \frac{t_s}{ft_s + (1-f)\frac{t_s}{p}} = \frac{p}{1 + (p-1)f} \quad [2]$$

where  $p$  is the number of processors used in the parallel algorithm,  $t_s$  and  $t_p$  are the respective running times for sequential and parallel runs of the algorithms, and  $f$  is the fraction of the work in a program that must be run sequentially. The second and third expressions are known as Amdahl's law. Ideally, the speedup factor will represent a number of factors  $S_p = p$ , but this is often not the case due to various factors including non-parallelizable segments of a job or inter-process communication overhead. If  $S_p$  is small relative to  $p$  or even approaches 1, then the parallel algorithm should either be improved or not used. Should  $S_p$  ever exceed  $p$ , this is known as *super-linear speedup*. A super-linear speedup factor is often the product of the use of a runtime that resulted from a sub-optimal sequential algorithm.

*Efficiency*

Parallel algorithms can also be measured in terms of how well the time to run a program is used by all of the processors. Efficiency measures how well the processors are used (what percentage of the runtime the processors spend working). Efficiency is found by:

$$E = \frac{t_s}{t_p * p} = \frac{S(p)}{p} \quad [3]$$

where  $t_s$ ,  $t_p$ , and  $S(p)$  are defined as in [2].

**Pseudo-Code**

It should be noted that the following pseudo-code should not be considered functional code. Functional code will be attached in the report.

*Sequential Pseudo-Code*

In a sequential program, it can be assumed that the work of a program is run as though it were only being run on a single processor. Thus operations are performed one after another, with no apparent re-ordering or concurrency.

```
double createMandelBrotImage( size )
{
    run_time
    Image_Matrix

    startTimer()

    // compute the pixels/values of each row
    for( i = 0; i < size; i = i + 1 )
    {
        for( j = 0; j < size; j = j + 1 )
        {
            Image_Matrix[i][j] = calcMandelPixel( i, j )
        }
    }

    run_time = stopTimer()

    createMandelbrotImage( Image_Matrix, size )

    return run_time
}
```

*Parallel Psuedo-Code*

In general, parallel programs have a *master* processor to direct the work and multiple *slave* processors to do most of the actual work. The master primarily performs communication and status checking operations, while the slaves primarily perform computations and only minimal sending of the results to the master.

*Static Job Allocation*

In a static job allocation algorithm each slave is allocated a set amount of work. In this algorithm, the slaves know which rows to compute based on their own number and the number of slaves. The slaves send the computation results to the master compiles the data.

Master:

```
double staticMaster( size )
{
    run_time
    Image_Matrix
    slave_id
    rows_left_to_compute = size
    row_to_receive

    startTimer()

    while( rows_left_to_compute > 0 )
    {
        receive( row_to_receive, slave_id )

        receive( Image_Matrix[row_to_receive], slave_id )

        rows_left_to_compute = rows_left_to_compute - 1
    }

    run_time = stopTimer()
}
```

Slave:

```
void staticSlave( size, slave_id, num_slaves )
{
    row_buffer
    current_row

    current_row = slave_id - 1

    while( current_row < size )
    {
        for( j = 0; j < size; j = j + 1 )
        {
            row_buffer[j] = calcMandelPixel( current_row, j )
        }

        send( row_buffer, master, slave_id )

        current_row = current_row + num_slaves
    }
}
```

### *Dynamic Job Allocation*

With dynamic job allocation, slaves are given a segment of work, and are to report for more work once their segment is complete. In my scheme, work is apportioned to slaves in blocks of 5 rows. Again, the master does administrative work, compiling data and delegating new work.

Master:

```
double dynamicMaster( size, num_slaves )
{
    run_time
    Image_Matrix
    slave_id
    next_row_to_compute
    num_received = 0
    row_recieved

    next_row_to_compute = ( num_slaves - 1 ) * size

    startTimer()

    while( num_received < size )
    {
        receive( row_received, slave_id )
    }
}
```

```

    for( k = 0; k < size; k = k + 1 )
    {
        receive( Image_Matrix[row_received + k], slave_id )
    }

    send( next_row_to_compute, slave_id )

    next_row_to_compute = next_row_to_compute + 10
}

run_time = stopTimer()
}

```

**Slave:**

```

void staticSlave( size, slave_id )
{
    row_buffer
    current_row

    current_row = ( slave_id - 1 ) * 10

    while( current_row < size )
    {
        for( i = 0; i < 5; i = i + 1 )
        {
            for( j = 0; j < size; j = j + 1 )
            {
                row_buffers[i][j] = calcMandelPixel(current_row + i, j )
            }
        }

        send( current_row, master, slave_id )

        for( i = 0; i < 5; i = i + 1 )
        {
            send( row_buffers[i], master, slave_id )
        }

        receive( current_row, master )
    }
}

```

## Output Data Summaries

The following tables and figures are representative of the output of the trial runs of the Mandelbrot set computation runs, listed by appropriate category. This is pretty boring and monotonous, but accurate, so if Devyani wanted to skip it, she could and still be confident that all the necessary items were here.

Computing Nodes Used in Experiment							
0:	compute-2-14.local	5:	compute-3-0.local	10:	compute-1-21.local	15:	compute-3-11.local
1:	compute-2-14.local	6:	compute-2-22.local	11:	compute-41-10.local	16:	compute-3-11.local
2:	compute-2-14.local	7:	compute-2-22.local	12:	compute-3-5.local	17:	compute-3-11.local
3:	compute-3-0.local	8:	compute-2-22.local	13:	compute-3-11.local	18:	compute-2-21.local
4:	compute-3-0.local	9:	compute-41-7.local	14:	compute-3-11.local	19:	compute-2-21.local

Table 1: The processor nodes used for the Mandelbrot computation runs. If a run used  $n$  processors, then processors  $n$  through  $n - 1$  were used.




Figure 1: a 10,000 x 10,000 pixel Mandelbrot image in grayscale produced by a sequential algorithm. Note: this image was converted from .pgm format to .jpeg format for Windows compatibility.




Figure 2: a 10,000 x 10,000 pixel Mandelbrot image in grayscale produced by a static job allocation parallel algorithm. Note: this image was converted from .pgm format to .jpeg format for Windows compatibility.




Figure 3: a 10,000 x 10,000 pixel Mandelbrot image in grayscale produced by a dynamic job allocation parallel algorithm. Note: this image was converted from .pgm format to .jpeg format for Windows compatibility.






Figure 4: a 3,000 x 3,000 pixel Mandelbrot image in color produced by a sequential computation run. A smaller image was used and a screen shot was taken in order to get this due to the limitations of my personal laptop and Windows. A real, 10,000 x 10,000 .ppm color image will be included with the other submitted files.




Figure 5: The trend of the time required to compute images of various sizes sequentially.

Image Size (pixels)	500	1000	1500	2000	2500	3000	3500	4000	4500	5000
Running Time (s)	0.147	0.562	1.100	1.909	2.975	4.274	5.755	7.484	9.562	11.791
Image Size (pixels)	5500	6000	6500	7000	7500	8000	8500	9000	9500	10000
Running Time (s)	14.132	16.869	19.675	22.782	26.287	29.815	33.754	37.812	42.082	46.507

Table 2: The benchmark average job computation times for computing Mandelbrot images of various sizes sequentially. These will act as the baselines for the later computations of speedup and efficiency.

Static Work Allocation Running Times										
Image Size (Pixels)	Number of Slaves									
	1	3	5	7	9	11	13	15	17	19
500	0.159	0.108	0.063	0.048	0.037	0.032	0.029	0.026	0.145	0.023
1000	0.501	0.244	0.176	0.194	0.126	0.257	0.248	0.222	0.239	0.338
1500	1.128	0.501	0.284	0.320	0.237	0.269	0.249	0.288	0.241	0.465
2000	1.907	0.795	0.479	0.441	0.327	0.344	0.322	0.412	0.302	0.288
2500	2.992	1.102	0.691	0.613	0.565	0.583	0.595	0.376	0.310	0.492
3000	4.283	1.503	0.932	0.727	1.025	0.737	0.523	0.489	0.423	0.553
3500	5.788	2.009	1.246	0.974	0.874	0.727	0.623	0.624	0.803	0.674
4000	7.523	2.580	1.592	1.397	1.006	0.973	0.966	0.752	0.632	0.638
4500	9.520	3.282	2.035	1.543	1.405	1.110	1.176	0.870	0.884	0.689
5000	11.693	3.999	2.416	2.025	1.584	1.347	1.365	1.104	1.066	0.841
5500	14.140	4.808	2.964	2.180	2.182	1.616	1.428	1.230	1.119	0.996
6000	16.800	5.762	3.474	2.802	2.276	1.922	1.613	1.409	1.316	1.117
6500	19.650	6.666	4.026	3.076	2.687	2.132	1.916	1.614	1.591	1.397
7000	22.791	7.704	4.708	3.627	3.058	2.525	2.263	2.098	1.681	1.547
7500	26.299	8.876	5.334	4.063	3.597	2.891	2.542	2.238	1.974	1.697
8000	29.755	10.068	6.045	4.811	4.094	3.356	2.820	2.479	2.281	1.828
8500	33.740	11.340	6.821	5.288	4.619	3.680	3.169	2.747	2.565	2.118
9000	37.541	12.704	7.785	5.848	4.996	4.223	3.502	3.564	2.802	2.644
9500	41.899	14.146	8.523	6.399	5.735	4.782	4.177	3.453	3.076	2.875
10000	46.626	15.725	9.478	7.211	6.107	5.068	4.214	3.915	3.418	3.126

Table 3: The average run times for static work allocation computation times by number of slaves used and image size.

Dynamic Work Allocation Running Times										
Image Size (Pixels)	Number of Slaves									
	1	3	5	7	9	11	13	15	17	19
500	0.147	0.074	0.101	0.227	0.049	0.085	0.046	0.027	0.079	0.063
1000	0.513	0.213	0.307	0.259	0.103	0.153	0.121	0.136	0.295	0.309
1500	1.093	0.417	0.450	0.298	0.238	0.193	0.190	0.215	0.329	0.401
2000	1.912	0.680	0.606	0.380	0.377	0.282	0.265	0.254	0.410	0.621
2500	2.969	1.034	1.287	0.586	0.430	0.375	0.378	0.344	0.340	1.078
3000	4.220	1.475	1.306	0.802	0.713	0.641	0.705	0.531	0.609	1.934
3500	5.753	1.944	1.680	1.159	1.466	1.118	0.905	1.482	0.719	2.420
4000	7.553	2.511	2.019	1.531	2.306	1.558	1.664	1.416	1.417	3.328
4500	9.482	3.146	2.239	1.909	3.238	2.568	2.261	1.695	1.959	3.982
5000	11.745	3.918	2.639	2.501	3.732	2.821	3.197	2.708	2.927	4.122
5500	14.224	4.684	3.393	4.065	4.645	4.093	3.744	2.911	2.504	5.528
6000	16.777	5.517	4.541	4.542	5.423	4.748	4.656	4.109	3.706	6.328
6500	19.752	6.500	4.263	5.717	5.800	4.950	5.053	4.466	5.184	8.222
7000	22.683	7.539	4.973	6.144	7.016	5.743	6.049	5.446	5.070	9.867
7500	26.299	10.590	5.860	7.087	6.815	7.720	6.457	5.884	6.262	9.255
8000	29.861	10.590	6.378	9.955	8.018	7.512	6.599	6.626	6.135	10.773
8500	33.679	12.442	7.406	11.580	8.574	8.144	7.566	6.809	7.394	11.023
9000	37.651	13.166	8.121	11.613	9.900	9.404	8.525	8.073	8.015	12.770
9500	41.916	14.293	9.835	10.817	10.917	13.340	10.729	8.980	9.795	15.218
10000	46.735	15.855	9.925	11.457	12.448	11.035	12.430	9.845	10.568	15.861

Table 4: The average run times for dynamic work allocation computation times by number of slaves used and image size.

## Speedup and Efficiency

Figures and tables representative of the Speedup and Efficiency results are displayed in table and graphical format in this section. It should be noted that due to ineffective message passing on the grid, some results are not as expected, the dynamic work allocation results in particular.

However, speedup of 8 was achieved for each job type after reaching an appropriate number of processing nodes/slaves.




Figure 6: A graph representing the speedup of the static Mandelbrot computation run averages.

Static Job Allocation Speedups										
	Number of Slaves									
Image Size (Pixels)	1	3	5	7	9	11	13	15	17	19
500	0.925	1.362	2.327	3.041	3.972	4.529	5.038	5.560	1.010	6.471
1000	1.123	2.306	3.194	2.897	4.476	2.191	2.265	2.528	2.349	1.662
1500	0.975	2.195	3.877	3.438	4.648	4.082	4.411	3.818	4.572	2.364
2000	1.001	2.401	3.984	4.333	5.846	5.542	5.932	4.630	6.323	6.634
2500	0.994	2.701	4.309	4.855	5.267	5.102	4.997	7.907	9.588	6.052
3000	0.998	2.844	4.584	5.879	4.171	5.803	8.176	8.736	10.095	7.725
3500	0.994	2.865	4.620	5.910	6.585	7.914	9.233	9.228	7.171	8.536
4000	0.995	2.901	4.702	5.355	7.436	7.690	7.747	9.946	11.839	11.725
4500	1.004	2.914	4.698	6.197	6.805	8.618	8.134	10.988	10.821	13.881
5000	1.008	2.949	4.880	5.822	7.446	8.754	8.640	10.678	11.058	14.021
5500	0.999	2.939	4.768	6.483	6.477	8.747	9.898	11.490	12.632	14.185
6000	1.004	2.928	4.855	6.021	7.413	8.776	10.459	11.971	12.818	15.105
6500	1.001	2.951	4.887	6.396	7.323	9.228	10.267	12.187	12.365	14.086
7000	1.000	2.957	4.839	6.282	7.451	9.021	10.065	10.857	13.551	14.729
7500	1.000	2.962	4.928	6.469	7.308	9.093	10.342	11.745	13.316	15.491
8000	1.002	2.961	4.932	6.197	7.282	8.884	10.573	12.027	13.068	16.315
8500	1.000	2.976	4.948	6.383	7.308	9.171	10.652	12.287	13.158	15.939
9000	1.007	2.976	4.857	6.465	7.569	8.953	10.798	10.608	13.494	14.301
9500	1.004	2.975	4.937	6.577	7.338	8.800	10.075	12.189	13.679	14.639
10000	0.997	2.958	4.907	6.449	7.615	9.176	11.036	11.880	13.608	14.880

Table 5: A table representing the speedup of the static Mandelbrot computation run averages.




Figure 7: A graph representing the speedup of the dynamic Mandelbrot computation run averages. Do note that the required speedup of 8 was achieved, if not by the majority of runs. I suspect this lack of speedup is due to the volatility of the grid communications.

Dynamic Job Allocation Speedups										
Image Size (Pixels)	Number of Slaves									
	1	3	5	7	9	11	13	15	17	19
500	1.002	1.979	1.449	0.647	3.023	1.736	3.208	5.508	1.848	2.320
1000	1.096	2.644	1.830	2.172	5.460	3.669	4.629	4.125	1.903	1.818
1500	1.006	2.635	2.442	3.693	4.621	5.688	5.798	5.125	3.348	2.746
2000	0.998	2.808	3.149	5.023	5.068	6.773	7.200	7.509	4.659	3.072
2500	1.002	2.878	2.313	5.073	6.917	7.933	7.867	8.637	8.751	2.761
3000	1.013	2.898	3.273	5.330	5.995	6.669	6.063	8.047	7.021	2.210
3500	1.000	2.961	3.427	4.967	3.926	5.149	6.362	3.884	8.004	2.378
4000	0.991	2.981	3.706	4.887	3.246	4.802	4.499	5.286	5.282	2.249
4500	1.008	3.039	4.271	5.009	2.953	3.723	4.229	5.640	4.882	2.401
5000	1.004	3.010	4.469	4.715	3.159	4.179	3.688	4.354	4.029	2.861
5500	0.994	3.017	4.165	3.477	3.042	3.453	3.775	4.854	5.644	2.557
6000	1.005	3.058	3.714	3.714	3.110	3.552	3.623	4.105	4.551	2.666
6500	0.996	3.027	4.615	3.442	3.392	3.975	3.894	4.406	3.795	2.393
7000	1.004	3.022	4.581	3.708	3.247	3.967	3.766	4.183	4.493	2.309
7500	1.000	2.482	4.486	3.709	3.858	3.405	4.071	4.467	4.198	2.840
8000	0.998	2.815	4.675	2.995	3.718	3.969	4.518	4.499	4.860	2.767
8500	1.002	2.713	4.557	2.915	3.937	4.145	4.461	4.957	4.565	3.062
9000	1.004	2.872	4.656	3.256	3.820	4.021	4.435	4.684	4.718	2.961
9500	1.004	2.944	4.279	3.890	3.855	3.154	3.922	4.686	4.296	2.765
10000	0.995	2.933	4.686	4.059	3.736	4.214	3.742	4.724	4.401	2.932

Table 6: A table representing the speedup of the dynamic Mandelbrot computation run averages.






Figure 8: A graph representing the efficiency of the static Mandelbrot computation run averages.

Static Job Allocation Efficiency										
Image Size (Pixels)	Number of Slaves									
	1	3	5	7	9	11	13	15	17	19
500	0.462	0.340	0.388	0.380	0.397	0.377	0.360	0.347	0.056	0.324
1000	0.561	0.577	0.532	0.362	0.448	0.183	0.162	0.158	0.131	0.083
1500	0.488	0.549	0.646	0.430	0.465	0.340	0.315	0.239	0.254	0.118
2000	0.501	0.600	0.664	0.542	0.585	0.462	0.424	0.289	0.351	0.332
2500	0.497	0.675	0.718	0.607	0.527	0.425	0.357	0.494	0.533	0.303
3000	0.499	0.711	0.764	0.735	0.417	0.484	0.584	0.546	0.561	0.386
3500	0.497	0.716	0.770	0.739	0.658	0.660	0.659	0.577	0.398	0.427
4000	0.497	0.725	0.784	0.669	0.744	0.641	0.553	0.622	0.658	0.586
4500	0.502	0.728	0.783	0.775	0.681	0.718	0.581	0.687	0.601	0.694
5000	0.504	0.737	0.813	0.728	0.745	0.730	0.617	0.667	0.614	0.701
5500	0.500	0.735	0.795	0.810	0.648	0.729	0.707	0.718	0.702	0.709
6000	0.502	0.732	0.809	0.753	0.741	0.731	0.747	0.748	0.712	0.755
6500	0.501	0.738	0.814	0.799	0.732	0.769	0.733	0.762	0.687	0.704
7000	0.500	0.739	0.807	0.785	0.745	0.752	0.719	0.679	0.753	0.736
7500	0.500	0.740	0.821	0.809	0.731	0.758	0.739	0.734	0.740	0.775
8000	0.501	0.740	0.822	0.775	0.728	0.740	0.755	0.752	0.726	0.816
8500	0.500	0.744	0.825	0.798	0.731	0.764	0.761	0.768	0.731	0.797
9000	0.504	0.744	0.810	0.808	0.757	0.746	0.771	0.663	0.750	0.715
9500	0.502	0.744	0.823	0.822	0.734	0.733	0.720	0.762	0.760	0.732
10000	0.499	0.739	0.818	0.806	0.762	0.765	0.788	0.743	0.756	0.744

Table 7: A table representing the efficiency of the static Mandelbrot computation run averages.




Figure 9: A graph representing the efficiency of the average dynamic Mandelbrot computation runs.

Dynamic Job Allocation Efficiency										
Image Size (Pixels)	Number of Slaves									
	1	3	5	7	9	11	13	15	17	19
500	0.501	0.495	0.242	0.081	0.302	0.145	0.229	0.344	0.103	0.116
1000	0.548	0.661	0.305	0.271	0.546	0.306	0.331	0.258	0.106	0.091
1500	0.503	0.659	0.407	0.462	0.462	0.474	0.414	0.320	0.186	0.137
2000	0.499	0.702	0.525	0.628	0.507	0.564	0.514	0.469	0.259	0.154
2500	0.501	0.720	0.385	0.634	0.692	0.661	0.562	0.540	0.486	0.138
3000	0.506	0.725	0.545	0.666	0.599	0.556	0.433	0.503	0.390	0.111
3500	0.500	0.740	0.571	0.621	0.393	0.429	0.454	0.243	0.445	0.119
4000	0.495	0.745	0.618	0.611	0.325	0.400	0.321	0.330	0.293	0.112
4500	0.504	0.760	0.712	0.626	0.295	0.310	0.302	0.352	0.271	0.120
5000	0.502	0.752	0.745	0.589	0.316	0.348	0.263	0.272	0.224	0.143
5500	0.497	0.754	0.694	0.435	0.304	0.288	0.270	0.303	0.314	0.128
6000	0.503	0.764	0.619	0.464	0.311	0.296	0.259	0.257	0.253	0.133
6500	0.498	0.757	0.769	0.430	0.339	0.331	0.278	0.275	0.211	0.120
7000	0.502	0.755	0.763	0.463	0.325	0.331	0.269	0.261	0.250	0.115
7500	0.500	0.621	0.748	0.464	0.386	0.284	0.291	0.279	0.233	0.142
8000	0.499	0.704	0.779	0.374	0.372	0.331	0.323	0.281	0.270	0.138
8500	0.501	0.678	0.760	0.364	0.394	0.345	0.319	0.310	0.254	0.153
9000	0.502	0.718	0.776	0.407	0.382	0.335	0.317	0.293	0.262	0.148
9500	0.502	0.736	0.713	0.486	0.385	0.263	0.280	0.293	0.239	0.138
10000	0.498	0.733	0.781	0.507	0.374	0.351	0.267	0.295	0.244	0.147

Table 8: A table representing the efficiency of the average dynamic Mandelbrot computation runs.

## Issues

The primary issues with this experiment were associated with the grid. Not only was the grid experiencing failures during the time this assignment, but the very nature of the grid led to inconsistencies in the data. This led to inconsistent job performance due to the varying levels of traffic on the grid, and the variety of hardware included in the network also led to inconsistent results. Quite frankly, the grid made this assignment awful, considering the poor performance of the grid that would not grant enough TCP connections or speedy message passing, bad etiquette from others, and long waiting periods (I had a job wait from 11:30am to 2:00 am to begin working at one point). Devyani and Dr. Harris should not take this personally, this section is just for identifying any hindrances to success on the project.

```

thenriod@hnode:~
thenriod@ubuntu: ~/Desktop/CS479/Project_1  thenriod@hnode:~
@compute-4-1.loc      4
 24709 0.56152 dynamic14 jblocher    r    03/09/2014 23:44:24 BrainLab.owner
@compute-41-1.lo     14
 24986 0.50935 Static    ahesson    r    03/09/2014 23:49:39 NSF.owner@comp
ute-3-7.local        2
 24748 0.60500 LS-DYNA-DA distratii  qw    03/09/2014 21:24:32
21725 0.58761 Mandelbrot thenriod    qw    03/09/2014 11:23:43
21934 0.58761 dynamic   davidfrank qw    03/09/2014 14:48:11
21952 0.58761 static    davidfrank qw    03/09/2014 15:33:38
24700 0.58761 static20   jblocher    qw    03/09/2014 19:21:19
24712 0.58761 dynamic20  jblocher    qw    03/09/2014 19:33:02
24764 0.58761 dynamicb   davidfrank  qw    03/09/2014 22:12:38
24766 0.58761 staticb    davidfrank  qw    03/09/2014 22:17:12
24696 0.55283 static12   jblocher    qw    03/09/2014 19:21:19
24708 0.55283 dynamic12  jblocher    qw    03/09/2014 19:33:02
24695 0.54413 static10   jblocher    qw    03/09/2014 19:21:19
24707 0.54413 dynamic10  jblocher    qw    03/09/2014 19:33:02
24694 0.53543 static8    jblocher    qw    03/09/2014 19:21:19
  
```

Figure 10: An example illustrating the difficulty of working on the grid in its current state. Note the drastic amount of time a job spent in the queue before getting to run. The terminal times are in 24:00 time, so this figure displays a >12 hour wait time.