

# CS 677: Assignment 5

Terence Henrion

April 20, 2014

## **Abstract**

In this assignment, various greedy algorithm problems and related ones are presented.

1. A search engine company needs to do a significant amount of computation every time it recompiles its index. For this task, the company has a single large supercomputer, and an unlimited supply of high-end PCs.

They have broken the overall computation into  $n$  distinct jobs, labeled  $J_1, J_2, \dots, J_n$ , which can be performed completely independently of one another. Each job consists of two stages: first it needs to be preprocessed on the supercomputer, and then it needs to be finished on one of the PCs. Lets say that job  $J_i$  needs  $p_i$  seconds of time on the supercomputer, followed by  $f_i$  seconds of time on a PC.

Since there are at least  $n$  PCs available on the premises, the finishing of the jobs can be performed fully in parallel all the jobs can be processed at the same time. However, the supercomputer can only work on a single job at a time, so the system managers need to work out an order in which to feed the jobs to the supercomputer. As soon as the first job in order is done on the supercomputer, it can be handed off to a PC for finishing; at that point in time a second job can be fed to the supercomputer; when the second job is done on the supercomputer, it can proceed to a PC regardless of whether or not the first job is done (since the PCs work in parallel); and so on.

Lets say that a *schedule* is an ordering of the jobs for the supercomputer, and the *completion time* of the schedule is the earliest time at which all jobs will have finished processing on the PCs. This is an important quantity to minimize, since it determines how rapidly El Goog can generate a new index.

Give a polynomial-time algorithm that finds a schedule with as small a completion time as possible.

*Note:* to prove that your greedy strategy yields the optimal solution, you have to prove that the problem has the greedy-choice property.

**Solution:** Intuition would tell us that we should try and achieve maximum parallelism in order to have a minimal completion time. In order to achieve this, we need to reduce the bottleneck effect that happens with the super-computer pre-processing and minimize the amount of time jobs spend waiting without being processed. Secondly, we need to order jobs such that jobs that need longer PC processing time get to start earlier, allowing the longer jobs to get some processing time in, leaving their remaining time the same or less than the jobs that need less PC processing time.

A greedy algorithm that would solve the problem in polynomial time would choose jobs with the shortest pre-processing time first, and in the event of a tie for pre-processing time, choose the job with the longest finishing time second. In order to accomplish this, we can simply sort the jobs using an  $n \lg n$  sorting algorithm, and feed the jobs into the system with the sorted order.

*The Algorithm*

**Algorithm** INDEX-REBUILD( $J(0 \dots n - 1)$ )

(\* Input: A list/array of jobs for index recompilation \*)

1. MODIFIED-RANDOMIZED-QUICKSORT( $J, 0, n - 1$ )
2. (\* This is the same as regular quicksort except for a different way of making comparisons, listed below \*)
3. **for**  $i \leftarrow 0$  **to**  $n - 1$
4.     feed a job to the super-computer

5. wait for the job to finish pre-processing
6. send the job to a PC after pre-processing to finish
7. **stop** (\* After jobs complete at the PCs, the index is recompiled. \*)

**Algorithm** *COMPARE-JOBS*( $A, B$ )

(\* Input: Two jobs to be re-compiled \*)

1. (\* Returns true if A is greater than (should be pre-processed after) B \*)
2. **if**  $A.p > B.p$
3.     **then**
4.         **return true**
5.     **else**
6.         **if**  $A.p < B.p$
7.             **then**
8.                 **return false**
9.             **else**
10.                 **if**  $A.f < B.f$
11.                     **then**
12.                         **return false**
13.                     **else**
14.                         **return true**

*Greedy Choice*

This algorithm is greedy because it makes the local optimum choice of selecting the shortest pre-processing time job first, and then the tandem choice of selecting the job with the longest finishing time if pre-processing times are equal.

*Greedy Choice Property*

To first prove the greedy choice of shortest pre-processing time first, consider the completion times  $c_i$  for each job, where  $c_i = \sum_{j=1}^n p_j$  since, clearly, the pre-processing completion times of each job depend on the jobs that run/complete before them.

Suppose that in the optimal solution the jobs run in some order where the  $k^{th}$  job has the shortest job time, call it  $p_k$ , and is not first job in the sequence, then the shortest job's completion time will be  $\sum_{i=1}^{k-1} p_i + p_k$ .

Consider also the first job in the sequence and any other job that will complete before the  $k^{th}$  job, whose completion times will be  $p_1$  and  $\sum_{i=1}^{k-1} p_i$  (or  $p_1 + \sum_{i=2}^{k-1} p_i$ ), respectively.

Suppose now, that the first job and the  $k^{th}$  job were swapped in the sequence. Note that the any job  $k+1$  and beyond will not have its completion time affected by this, so these will not be addressed further. Not the completion time for the  $k^{th}$  will be just  $p_k$ , the completion time for all jobs 2 to  $k-1$  is  $p_k + \sum_{i=2}^{k-1} p_i$ , and the completion time for job 1 is  $p_k + \sum_{i=2}^{k-1} p_i + p_1$  (or  $p_k + \sum_{i=1}^{k-1} p_i$ ).

Consider now the differences in the original completion times for a job in the  $i^{th}$  position in the sequence to those after the swap. The first job run will complete  $(p_1 - p_k)$  time units faster (which is desirable because it was stated that  $p_k < p_1$ ), the completion times of jobs 2 to  $k-1$  will be  $(p_1 + \sum_{i=2}^{k-1} p_i) - (p_k + \sum_{i=2}^{k-1} p_i) = p_1 - p_k$ , again, resulting in faster completion times, and finally the difference in completion time for job 1 will be  $(\sum_{i=1}^{k-1} p_i + p_k) - (p_k + \sum_{i=1}^{k-1} p_i) = 0$ , neither faster nor slower. Observe that by making the swap, all completion times except the  $k^{th}$  position one were better, and there was no disadvantage for the job in the  $k^{th}$  position! It is clearly advantageous to run the shortest running time job first. This will be true of any set of jobs, including the set of remaining jobs after a shortest job is chosen and removed from the original set.

The greedy choice of choosing the job with the shortest pre-processing time first has improved on what we had originally assumed to be the optimal solution (for the purposes of this proof, the above logic could be used to directly prove/derive the *actual* the optimal solution).

For the second greedy choice, where we select the job with the longest regular PC processing time  $f_i$  first in the case of a tie for pre-processing time  $p_i$  (or equivalently, choose the job with the shortest finishing time to go last). For this part, completion times are independent of one another since all jobs can be run in parallel, the only assumption is that no jobs will start the finishing process at the same time if pre-processing times are non-zero. In this case, we need to maximize for the earliest finishing time for each job. Since the previous choice already determined which the order of pre-processing jobs, this choice will only re-order jobs with equal pre-processing times.

Consider a set of  $m$  jobs, all with the same pre-processing time  $p$ , but differing finishing times  $f_i$ . Let the absolute finishing time of a job be called  $d_i$ , where  $d_i = (i - 1) * p + f_i$ .

Take a set (or pair) of any two jobs with different finishing times. Let the finishing time of the longer job be called  $f_l$  and the shorter one called  $f_s$ . If the jobs are ordered with the shorter job first, then its absolute finishing time is  $d_s = (1 - 1) * p + f_s = f_s$ , and the longer job's absolute finishing time will be  $d_l = (2 - 1) * p + f_l = p + f_l$ . The greedy choice, however, would take the longer job first, giving absolute finishing times of  $d_l = (1 - 1) * p + f_l = f_l$  and  $d_s = (2 - 1) * p + f_s = p + f_s$ , respectively. The difference between each of the absolute finishing times in each variation is just  $p$ , depending on whether or not the greedy choice was used. Now, either  $(f_l - f_s) > p$ , or it is not. If  $(f_l - f_s) > p$ , then  $d_l > d_s$  using either strategy, but  $d_l$  will be earlier/smaller using the greedy strategy ( $f_l < p + f_l$ ), meaning the greedy strategy is superior in that case. If  $(f_l - f_s) \leq p$ , then with either strategy, the greatest absolute finishing time will be the one corresponding to whichever job is run second. For the "optimal" decision this is  $d_l = p + f_l$ , and for the greedy method this is  $d_s = p + f_s$ . Clearly  $p + f_s < p + f_l$ , so once again the greedy method is superior. Thus, the greedy choice is superior to the "optimal" one.

This logic can be expanded to a set of any size by ordering the jobs by arranging them pair-wise in order to achieve the best pair-wise absolute finishing times, and repeating until no improvements can be made.

The entire greedy choice has been shown to be the same as or better than an optimal strategy.

#### *Optimal Substructure*

As is intuitive, the optimal algorithm for ordering  $n$  jobs will be the same algorithm to use to order  $n - 1$  jobs after the first one has been selected. This is just a formality to justify our use of a greedy strategy.

#### *Polynomial Time*

This algorithm is easily seen to be a polynomial time algorithm. It used a modified efficient sorting algorithm, like quick sort, where the modification is that each comparison is actually as many as 3 comparisons are made instead of one, which is still a  $O(n \lg n)$  sorting algorithm. After the sorting, we can feed each of the  $n$  jobs in to the computers into the supercomputer and then into a PC (2 operations for each job), for a cost of  $\Theta(2n) = \Theta(n)$  operations. The resulting algorithm will then be  $\Theta(n) + \Theta(n \lg n) = \Theta(n \lg n)$ , which is polynomial.

2. Suppose you have  $n$  video streams that need to be sent, one after another, over a communication link. Stream  $i$  consists of a total of  $b_i$  bits that need to be sent, at a constant rate, over a period of  $t_i$  seconds. You cannot send two streams at the same time, so you need to determine a schedule for the streams: an order in which to send them. Whichever order you choose, there cannot be any delays between the end of one stream and the start of the next. Suppose your schedule starts at time 0 (and therefore ends at time  $\sum_{i=1}^n t_i$ , whichever order you choose). We assume that all the values  $b_i$  and  $t_i$  are positive integers.

Now, because you are just one user, the link does not want you taking up too much bandwidth, so it imposes the following constraint, using a fixed parameter  $r$ :

(\*) For each natural number  $t > 0$ , the total number of bits you send over the time interval from 0 to  $t$  cannot exceed  $rt$ .

Note that this constraint is only imposed for time intervals that start at 0, not for time intervals that start at any other value.

We say that a schedule is *valid* if it satisfies the constraint (\*) imposed by the link.

**The problem.** Given a set of  $n$  streams, each specified by its number of bits  $b_i$  and its time duration  $t_i$ , as well as the link parameter  $r$ , determine whether there exists a valid schedule.

**Example.** Suppose we have  $n = 3$  streams, with

$$(b_1, t_1) = (2000, 1),$$

$$(b_2, t_2) = (6000, 2),$$

$$(b_3, t_3) = (2000, 1),$$

and suppose the link's parameter is  $r = 5000$ . Then, the schedule that runs the streams in the order 1, 2, 3 is valid, since the constraint (\*) is satisfied:

$t = 1$ : the whole first stream has been sent, and  $2000 < 5000 * 1$

$t = 2$ : half of the second stream has also been sent, and  $2000 + 3000 < 5000 * 2$ .

Similar calculations hold for  $t = 3$  and  $t = 4$ .

- (a) Consider the following claim:

There exists a valid schedule if and only if each stream  $i$  satisfies  $b_i \leq r * t_i$ .

Decide whether you think the claim is true or false, and give a proof of either the claim or its negation.

**Solution:** The claim is false. It is possible that a stream could violate  $b_i \leq r * t_i$  on its own, but be part of a schedule where the average  $b_{total}$  is  $\leq r * t_{total}$ , i.e.  $\sum_{i=1}^n b_i \leq r * \sum_{i=1}^n t_i$ . Consider a set of streams being defined by  $(b_1, t_1) = (1500, 1)$  and  $(b_2, t_2) = (500, 1)$ , with  $r = 1000$ . In this case, stream 1 does not satisfy the  $b \leq r * t$  requirement alone, since  $1500 > 1000 * 1$ . However, we could schedule the jobs so that the stream order is 1, 2, in which case we have  $b_{total} = 500 + 1500 = 2000$  and  $t_{total} = 1 + 1 = 2$ , giving  $2000 \leq 1000 * 2$ , which satisfies (\*) from the problem statement, and is therefore a valid schedule.

- (b) Give an algorithm that takes a set of  $n$  streams, each specified by its number of bits  $b_i$  and its time duration  $t_i$ , as well as the link parameter  $r$ , and determines whether there exists a valid schedule. The running time of your algorithm should be polynomial in  $n$ .

**Solution:** A greedy choice that would be good for selecting a schedule for this scenario would be to select the stream with the lowest “bit density” first, and continue doing so from the remaining set to determine the order. This problem exhibits an optimal sub-structure because finding a valid way to schedule  $n$  streams must also find a valid way to schedule  $n - 1$  streams so as to keep

$b_{tot} \leq r * t_{tot}$  for each set. The greedy choice is as good or better than the “optimal” algorithm because if an optimal algorithm were to schedule the streams in some valid order without the lowest “bit density” stream being first, making the greedy choice of scheduling the lowest bit density stream first would not alter the average bit density of the schedule, and would in fact result in a schedule that kept the bit density further below the threshold of the condition than would a different ordering.

The question says to give an algorithm that simply indicates if there exists a valid scheduling. This is almost trivial since all that is required is to total  $b_i$ s and the  $t_i$ s and check if  $b_{tot} \leq r * t_{tot}$ , and if that is satisfied then there exists a valid scheduling. It is of course polynomial since it requires  $2 * (n - 1)$  additions and 1 comparison, which is  $2n - 1 = \Theta(n)$ .

This is boring and not as good as actually determining the valid schedule. A good algorithm for determining the existence of a valid schedule *and* finding such a valid schedule in polynomial time would look like the following:

**Algorithm** *FIND-VALID-SCHEDULE*( $S(0 \dots n - 1), r$ )

(\* Input: A list or array of streams \*)

1. (\* An array of the “bit densities” of the streams and an identifier for determining the where in the order each stream goes \*)  $D(0 \dots n - 1)$
2.  $density_{tot} \leftarrow 0$
3. **for**  $i \leftarrow 0$  **to**  $n - 1$
4.     **do**
5.          $D[i] = S[i].b \ S[i].t$
6.         *RANDOMIZED-QUICKSORT*( $D, 0, n - 1$ )
7. (\* Sort the jobs by bit density \*)
8. **for**  $i \leftarrow 0$  **to**  $n - 1$
9.      $density_{tot} = density_{tot} + D[i]$
10.     **if**  $density_{tot} > r$
11.         **then**
12.             **return false** (\* There is not a valid schedule of the  $n$  streams \*)
13. **for**  $i \leftarrow 0$  **to**  $n - 1$
14.     Run stream  $D[i].identifier$
15. **return true** (\* There is a valid schedule of the  $n$  streams \*)

This is a  $\Theta(n * 2) + \Theta(n \lg n) + \Theta(n) = \Theta(n \lg n)$  algorithm.

3. Exercise 16.1-2 (page 422): Suppose that instead of always selecting first activity to finish, we instead select the last activity to start that is compatible with all previously selected activities. Describe how this approach is a greedy algorithm and prove that it yields an optimal solution.

**Solution:** This algorithm works similarly to the typical greedy algorithm for scheduling the maximum number of activities, just in reverse. That is, this algorithm starts at the “end” of the day (as opposed to the beginning of the day), and picks the activity that starts the latest (as opposed to the one that finishes the earliest), which leaves a sub-problem of selecting the next activity from the set of activities that all finish before the selected one starts (as opposed to leaving activities that all start after the recently selected activity finishes).

This algorithm is a greedy one because it chooses the local optimum choice at each step: the activity with the latest start time in the set of remaining compatible activities is chosen. Further, as we already know, the problem has an optimal sub-structure: the optimal solution to the problem relies on finding the optimal solution to the sub-problems, which arises after every selection of an activity, which, the algorithm asserts is optimal.

To show that this solution is optimal, consider the set that includes the most compatible activities using the specified algorithm, call it  $A_{max}$ , which, of course, is a sub-set of all considered activities  $S_k$ . Let  $a_i$  be the activity in  $A_{max}$  that has the latest starting time. Suppose we choose the activity with the latest starting time from  $S_k$ , call it  $a_l$ . There are two cases of this: either  $a_l$  is the same activity as  $a_i$ , or  $a_l$  is not  $a_i$ , and therefore has an earlier start time than  $a_i$ . The first case is trivial; if  $a_l$  is the same as  $a_i$ , then clearly the choice has not detracted from the optimal solution (because it was the same choice as the optimal solution), and produces an activity set the same size as the optimal solution (because it is the same). In the other case where, however, we could form a new set that must be the same or better than  $A_{max}$ , we could call it  $A'_{max}$ . We know that  $A'_{max}$  is the same or better than  $A_{max}$  because  $a_l$  is compatible with all of the same activities that  $a_i$  is (any activity that is compatible with  $a_i$  must finish before  $a_i$ , so since  $a_l$  starts later than  $a_i$ , clearly the same activities are compatible with  $a_l$ ).  $A'_{max}$  might even be better because the later start time might allow for more compatible options (it should be noted that this is not possible because one cannot do better than the theoretical optimum solution). In this case, the resulting set will be of the same size or greater than the optimal set.

In either case, the greedy choice led to a solution set ( $A_{max}$  or  $A'_{max}$ ) that was of the same size as the optimal solution, therefore the greedy algorithm is valid and finds the optimal solution.

Finally, after selecting an activity, due to the optimal sub-structure of the problem, we can repeat the above logic on smaller sets of activities until there are no activities remaining.

4. Exercise 16.3-1 (page 436): Explain why in the proof Lemma 16.2, if  $x.freq = b.freq$ , then we must have  $a.freq = b.freq = x.freq = y.freq$ .

**Solution:** Since  $x$  and  $b$  are defined as being the letters with the lowest frequency and the highest frequency, respectively, we can consider their frequencies as the endpoints of an interval of numbers, positive integers if we must be precise. All numbers  $j_i$  in the interval of counts must then satisfy  $x.freq \leq j_i \leq b.freq$ . Clearly, the only way for a number to be in this interval if  $x.freq = b.freq$  is for  $j_i = x.freq = b.freq$ . Thus, since  $a.freq$  and  $y.freq$  are in this interval, they must satisfy  $a.freq = b.freq = x.freq = y.freq$ .