

Attempting to Improve APSP on the GPU

Terence Henriod

Intro

- Graphs are an abstraction of binary relationships between entities and the cost or weights of those relationships
- Sometimes we want to know the best way to navigate between one entity and another

- Sometimes we are interested in optimizing cost or reward of traversal, such as in the case of a road map or a probabilistic model
- Maybe we want to know the best places to start from in order to disperse something (related to the largest clique problem)
- No matter what, we want to find our answer FAST

My Results In Short...

- I failed to improve upon the APSP algorithm in a significant way
- I have found ways that don't work or only work in non-target cases → maybe I have pruned the “search tree” of ways to improve
- I attempted the important things that I proposed and added some others

My Ideas for Improvement

1. **Solve subtrees of the recursion tree on separate, concurrent GPUs**
2. **Increase the β parameter to shorten the recursion tree**
3. **Increase the number of threads used in one of the existing kernels**
4. **Decrease the number of threads used (I'll explain)**
5. Use shared memory to execute *some* of the kernels concurrently
6. Try to improve on Volkov's SGEMM
7. Choose a different project (just kidding)

Solving Subtrees Independently

- This failed due to my own oversight
- Could've seen this by studying the math/regular expressions more closely; could've tried reordering the operations sooner rather than writing code to re-arrange it
- In short: impossible, one subtree is dependent on the solution of the other.

$$A^* : \mathbb{R}^{N \times N} = \text{APSP}(A : \mathbb{R}^{N \times N})$$

```

1  if  $N < \beta$ 
2      then  $A \leftarrow \text{FW}(A)$             $\triangleright$  Base case: perform iterative FW serially
3  else
4       $A = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}$ 
5       $A_{11} \leftarrow \text{APSP}(A_{11})$ 
6       $A_{12} \leftarrow A_{11}A_{12}$ 
7       $A_{21} \leftarrow A_{21}A_{11}$ 
8       $A_{22} \leftarrow A_{22} \oplus A_{21}A_{12}$ 
9       $A_{22} \leftarrow \text{APSP}(A_{22})$ 
10      $A_{21} \leftarrow A_{22}A_{21}$ 
11      $A_{12} \leftarrow A_{12}A_{22}$ 
12      $A_{11} \leftarrow A_{11} \oplus A_{12}A_{21}$ 

```

Can be Rearranged

Can be Rearranged

Fig. 3. Pseudocode for recursive in-place APSP.

DOH!

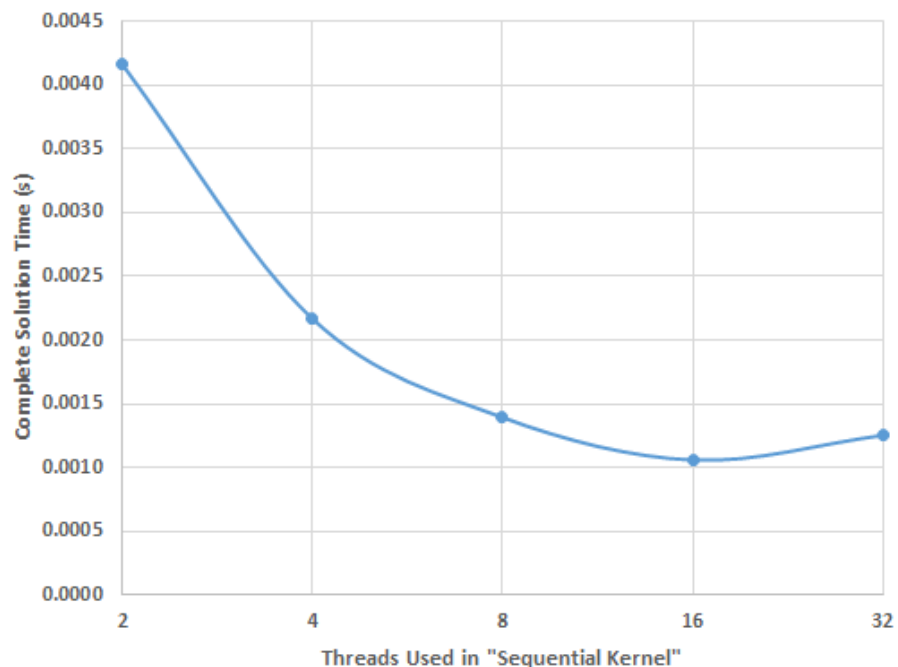
Increase the β Parameter

- Reduce recursion depth, reduce kernel launches, and maybe eliminate the need for a 3rd kernel!
- The kernels used:
 - if ($\text{block_width} < \beta$) \rightarrow “sequential” multiply
 - **else if ($\text{block_width} < 64$) \rightarrow “regular” multiply**
 - else \rightarrow “Volkov” multiply

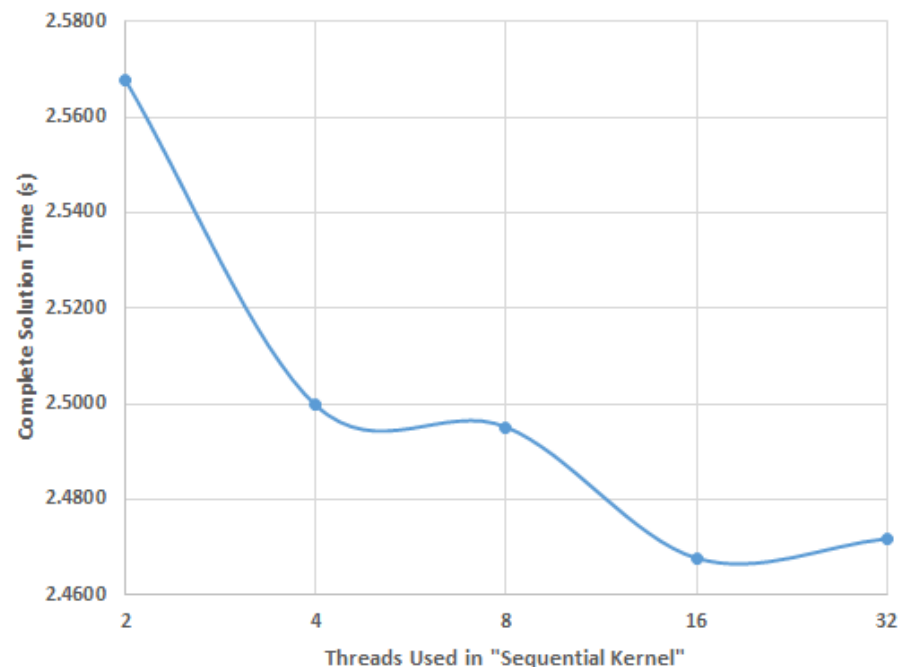
Increase the β Parameter / Increase the Number of Threads Used

- Attempt to reduce the number of kernel launches/have a faster kernel
- It's not very effective...
- Anecdotally: slightly better on a mobile GPU (but that wasn't the goal...)
- Actually a little worse than the 16 originally used on a 480

Execution of Original Code with Varied Threads in
"Sequential" FW Kernel for 256 x 256 Graph



Execution of Original Code with Varied Threads in
"Sequential" FW Kernel for 8192 x 8192 Graph



Increase the β Parameter/Use Fewer Threads

- Again, attempt to reduce the number of kernel launches, also reduce global memory accesses, and give threads more work
- Not always correct - lost correctness above 8 threads
- Achieved speedup on smaller inputs/vs. smaller numbers of threads (but that's not the goal...)

- The strategy:

- Fetch memory into shared memory to reduce global memory accesses
- Have threads do more work by having each thread compute results for 4 or 16 elements of the result rather than 1 (exploit Instruction Level Parallelism; see Volkov's work <http://www.cs.berkeley.edu/~volkov/volkov10-GTC.pdf>)
- Should be correct since in the original algorithm, each thread assessed one pair of neighbors independently at a time; I am simply simulating a larger number of threads doing the same thing.

- Assumed values for example analysis:
 - 4 x 4 thread block → 4 x 4 small block, 8 x 8 large block
 - 600 cycles (~150 GB/s) per global memory access
 - 6 (~1.7 TB/s) for shared memory accesses
 - reads/writes considered similar for simplicity (I know they are different)
- Original Method Memory Access
 - each thread (4^2 threads) performs 2 reads and 1 write 4 times
→ $4^2 * (2 + 1) * 4 = 192$ global accesses per call
 - 4 small blocks in a big block
 - 4 small blocks * 192 global accesses * 600 cycles = 460800 total cycles
- My method
 - each thread performs 4 global memory reads (start) and 4 writes (end)
 - each thread (4^2 threads) performs 8 shared memory reads and 4 writes 8 times → $4^2 * (8 + 4) * 8 = 1536$ shared memory accesses
 - 1 big block * ($1536 * 6 + 8 * 600$) = 14016 total cycles
- $14016 \ll 460800$


- In Big-Oh

- Original

- $4 \text{ tiles} * (t^2 * (2gr + 1 gw) * 4)$
- $= 32 * t^2 * gr + 16 t^2 * gw$
- $= O(28800 * t^2)$

- Mine: $1 \text{ tile} * (t^2 * (4gr + 4gw + 8sr + 4sw))$
- $= 4*t^2*gr + 4*t^2*gw + 8*t^2*sr + 4*t^2*sw$
- $= O(4860 * t^2)$

***If we stick to our 600 cycles for global and 6 cycles for shared assumption




Original Method

Example: The thread with
threadId.x == 1,
threadId.y == 1

x16 threads in both cases

Each operation is presumably
independent.



My Method

- Could it be that allocating a small array in registers in the “regular” multiplies is better than dynamically allocating shared memory in the “sequential” multiply?

Didn't Try: Execute *Some* Kernels Independently

- Without any machines with *integrated* GPUs available, this seemed like a lost cause
- With physically unified memory, using “pinned memory” techniques or true Unified Memory (UM) might have produced some improvement
- As it stands, the hit from accessing non-UM multiple times from a discrete GPU probably would have been detrimental

Didn't Try: Improve Volkov's SGEMM

- Volkov is a wiz at SGEMM - likely little room for improvement
- There seems to be intricate hand tuning involved - low probability of improving it anyway

```

__global__ void sgemmNN_MinPlus(
const float *A, int lda, const float *B, int ldb, float* C,
int ldc, int k, float beta ) {
|
const int inx = threadIdx.x;
const int iny = threadIdx.y;
const int ibx = blockIdx.x * 64;
const int iby = blockIdx.y * 16;
const int id = inx + iny*16;

A += ibx + id;
B += inx + __mul24( iby + iny, ldb );
C += ibx + id + __mul24( iby, ldc );

const float *Blast = B + k;

float c[16] = {FLOATINF, FLOATINF, FLOATINF, FLOATINF,
               FLOATINF, FLOATINF, FLOATINF, FLOATINF,
               FLOATINF, FLOATINF, FLOATINF, FLOATINF,
               FLOATINF, FLOATINF, FLOATINF, FLOATINF};

do {
    float a[4] = { A[0*lda], A[1*lda], A[2*lda], A[3*lda] };

    __shared__ float bs[16][17];
    bs[inx][iny] = B[0*ldb];
    bs[inx][iny+4] = B[4*ldb];
    bs[inx][iny+8] = B[8*ldb];
    bs[inx][iny+12] = B[12*ldb];
    __syncthreads();

```

```

A += 4*lda;
saxpy_MinPlus<16>( a[0], &bs[0][0], c );    a[0] = A[0*lda];
saxpy_MinPlus<16>( a[1], &bs[1][0], c );    a[1] = A[1*lda];
saxpy_MinPlus<16>( a[2], &bs[2][0], c );    a[2] = A[2*lda];
saxpy_MinPlus<16>( a[3], &bs[3][0], c );    a[3] = A[3*lda];

A += 4*lda;
saxpy_MinPlus<16>( a[0], &bs[4][0], c );    a[0] = A[0*lda];
saxpy_MinPlus<16>( a[1], &bs[5][0], c );    a[1] = A[1*lda];
saxpy_MinPlus<16>( a[2], &bs[6][0], c );    a[2] = A[2*lda];
saxpy_MinPlus<16>( a[3], &bs[7][0], c );    a[3] = A[3*lda];

A += 4*lda;
saxpy_MinPlus<16>( a[0], &bs[8][0], c );    a[0] = A[0*lda];
saxpy_MinPlus<16>( a[1], &bs[9][0], c );    a[1] = A[1*lda];
saxpy_MinPlus<16>( a[2], &bs[10][0], c );   a[2] = A[2*lda];
saxpy_MinPlus<16>( a[3], &bs[11][0], c );   a[3] = A[3*lda];


A += 4*lda;
saxpy_MinPlus<16>( a[0], &bs[12][0], c );   a[0] = A[0*lda];
saxpy_MinPlus<16>( a[1], &bs[13][0], c );   a[1] = A[1*lda];
saxpy_MinPlus<16>( a[2], &bs[14][0], c );   a[2] = A[2*lda];
saxpy_MinPlus<16>( a[3], &bs[15][0], c );   a[3] = A[3*lda];

B += 16;
__syncthreads();
} while( B < Blast );

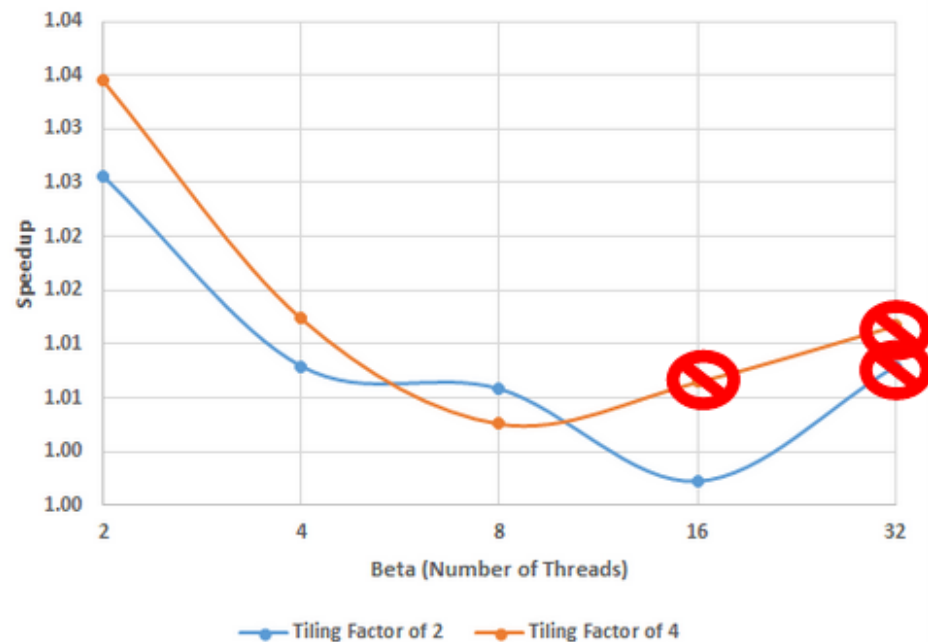
for( int i = 0; i < 16; ++i, C += ldc )
{
    C[0] = fminf(c[i],beta*C[0]);
}
}

```

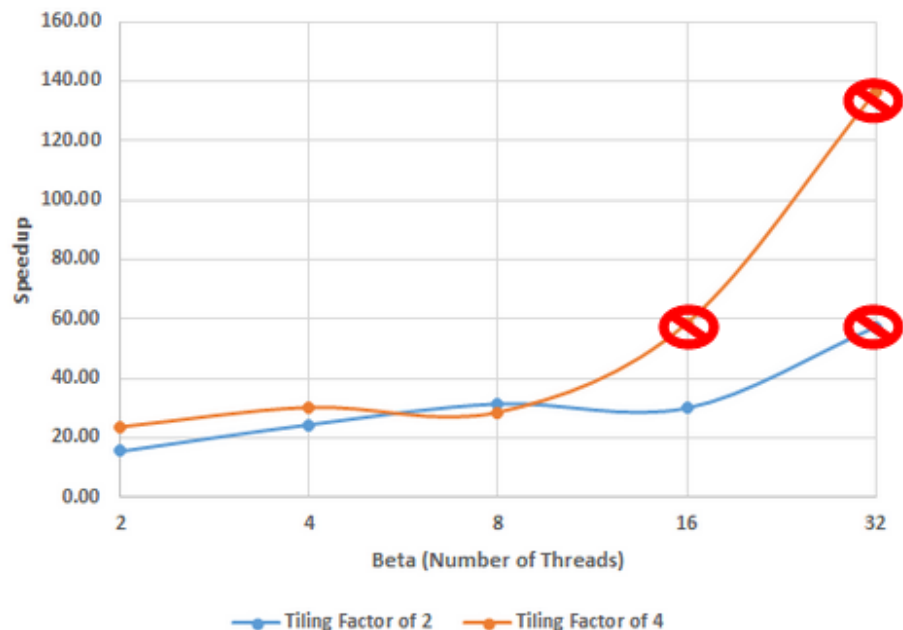
Speedup of the Increased Beta/Decreased Threads Computation over Original GPU Implementation for 256 x 256 Matrix



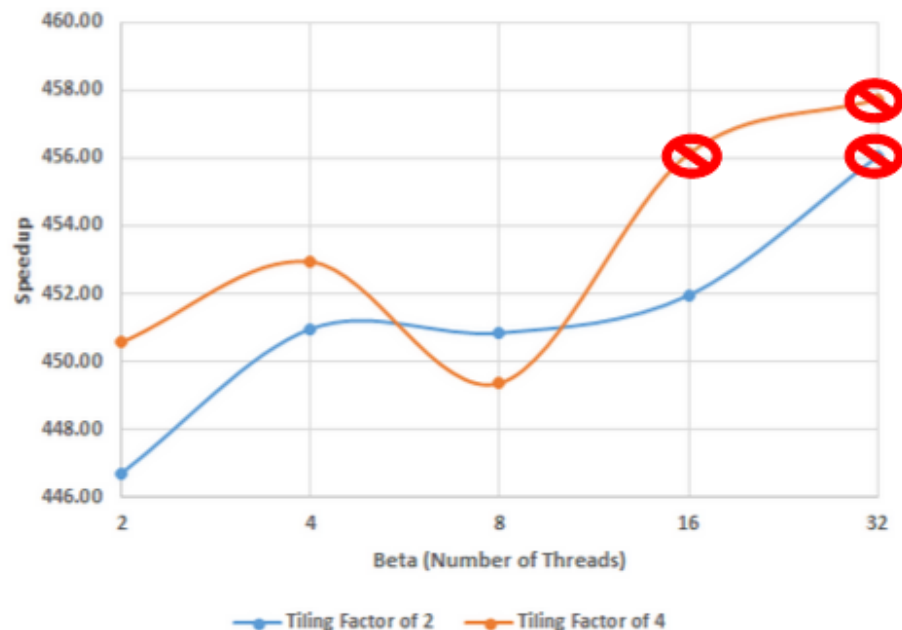
Speedup of the Increased Beta/Decreased Threads Computation over Original GPU Implementation for 8192 x 8192 Matrix



Speedup of the Increased Beta/Decreased Threads
Computation over Sequential Implementation for
256 x 256 Matrix



Speedup of the Increased Beta/Decreased Threads
Computation over Sequential Implementation for
8192 x 8192 Matrix



The Hardware

- CPU

- Intel Core i7-4790
- 15.6 GiB RAM
- 8 MB cache
- 3.60 GHz

- GPU

- GeForce GTX 480
- 15 SM; 448 Cores
- 1280 MB RAM
- 786 kB L2 cache
- 1.4 GHz

If I had another month...

- I'd allow myself a week to try and fix the increased β /fewer threads approach
- Spend 2 days (at most) trying the dual GPU with shared memory for concurrent kernels approach
- Add blocks to the “normal” multiply kernel
- Try out concurrent kernel launches where possible
- Tuning or improving Volkov's SGEMM

Thank You

Questions?